

Gymnasium (essentially gym)

- A project that provides an API for single agent RL environments
- Four key functions: `make`, `Env.reset()`, `Env.step()`, and `Env.render()`
- `Env` is a python class representing a markov decision process
- Allows creation of an initial state, transition/move to new states, and being able to visualize the environment.
- Wrapper can be used to augment/modify the environment

To initialize an environment, use the `make()` function. An example (derived from original documentation)

```
...  
import gymnasium as gym  
env = gym.make('CartPole-v1')  
...
```

Now, an `env` function is returned.

`pprint_registry()` shows all creatable environments

**Code from the original documentation; creates a single-episode run of lunar lander:**

```
import gymnasium as gym  
  
env = gym.make("LunarLander-v3", render_mode="human")  
observation, info = env.reset()  
  
episode_over = False  
while not episode_over:  
    action = env.action_space.sample() # agent policy that uses the observation and info  
    observation, reward, terminated, truncated, info = env.step(action)  
  
    episode_over = terminated or truncated  
  
env.close()
```

**End code.**

An environment is first created with `make()`. Render mode specifies how to make the environment. The specified `LunarLander` environment is made.

Using `Env.reset()` on the environment gets the first environment observation with some additional info.

We want to continue the agent environment loop till the environment ends, which we don't know when that will happen. `Episode_over` is a variable that determines when to stop the environment.

The agent performs an action in this environment; `Env.step()` is what executes the specified step. This action becomes random with `env.action_space.sample()`. This is like having a robot press a specified button on a controller to play a video game. From taking that action, there is a new environment and a reward for taking that action. (positive for doing something good, negative for doing something bad). This whole paragraph details a timestamp.

The environment may end after some steps (if the robot gets to a state where it can't continue, the robot must stop. In this case, `step()` returns `terminated`).

If we want the environment to end after X number of steps, the environment issues a truncated signal.

Receiving any end signal like this usually results in restarting the environment

`Action_space` and `observation_space` attributes for the environment. This is good for knowing the input and output of the environment. We might want to use `env.action_space.sample()` instead of an agent policy.

Most importantly, `Env.action_space` and `Env.observation_space` are instances of `Space`

There are a variety of spaces, such as a box, which has upper and lower limits of an n dimensional shape. There is `Text`, which describes a string space with min/max length, or a dictionary or tuple, which are self-explanatory.

To modify the environment:

Wrappers can do this (without having to alter the underlying code). Will allow code to be more modular and can also be changed

- When using `gymnasium.make()`, generally will be wrapped with `TimeLimit`, `OrderEnforcing`, and `PassiveEnvChecker`.

First need to make a base environment. Can then pass environment with (possibly optional) parameters to the wrapper's constructor:

When having a wrapped environment and you want to get the unwrapped environment underneath all layers of wrappers, use the `unwrapped` attribute. Unwrapped attributes on a base environment will just return itself.

There are a few wrappers. These include:

- TimeLimit, which issues a truncated signal if the program exceeds the maximum amount of timestamps.
- ClipAction, which clips any action passed to step such that it lies in the base environment's action space
- RescaleAction: applies an affine transformation to a particular action to linearly scale for a new low and high bound on the environment
- TimeAwareObservation: adds info about the index of timestep to observation.